

```
In [435]: from pmdarima import auto_arima
import pandas as pd
import matplotlib.pyplot as plt
import requests
import io
import seaborn as sns
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import adfuller
import statsmodels.api as sm
from statsmodels.tsa.seasonal import seasonal_decompose

# URL of the dataset
url = "https://fred.stlouisfed.org/graph/fredgraph.csv?id=ICSA"

# Fetching the data from the URL
response = requests.get(url)

# Reading the data into a DataFrame
df = pd.read_csv(io.StringIO(response.text))

print(df.head(10)) # Displays the first 10 rows
```

	DATE	ICSA
0	1967-01-07	208000
1	1967-01-14	207000
2	1967-01-21	217000
3	1967-01-28	204000
4	1967-02-04	216000
5	1967-02-11	229000
6	1967-02-18	229000
7	1967-02-25	242000
8	1967-03-04	310000
9	1967-03-11	241000

```
In [436]: df.shape
```

```
Out[436]: (2980, 2)
```

```
In [437]: # Convert 'DATE' column to datetime
print('DATE data type before conversion:')
print(df['DATE'].dtypes)
df['DATE'] = pd.to_datetime(df['DATE'])
print('DATE data type after conversion:')
print(df['DATE'].dtypes)
```

DATE data type before conversion:
object
DATE data type after conversion:
datetime64[ns]

```
In [438]: # Check for missing values
missing_values = df.isnull().sum()
# Check for duplicates
duplicates = df.duplicated().sum()

# Handling missing values (if any, here we just print them)
print(f"Missing values:\n{missing_values}")
print(f"Duplicate values:\n{duplicates}")
```

Missing values:
DATE 0
ICSA 0
dtype: int64
Duplicate values:
0


Here there are no missing or duplicate values

```
In [439]: print(f"Statistical Summary:")
print(df.describe())
```

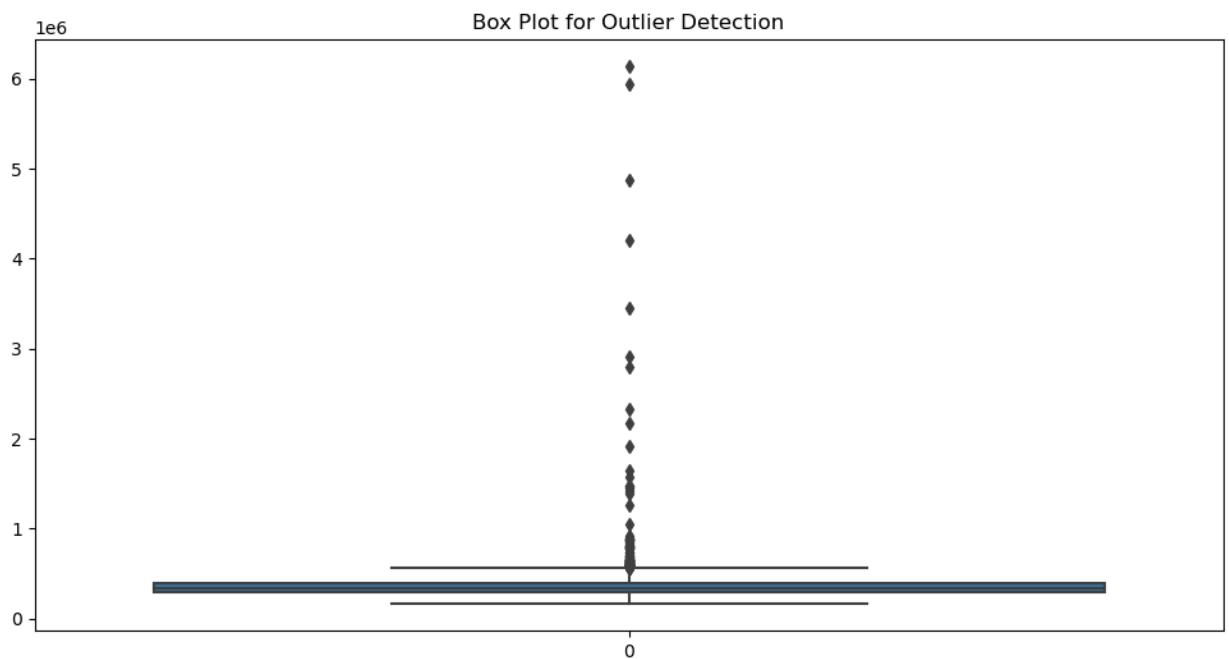
Statistical Summary:

	DATE	ICSA
count	2980	2.980000e+03
mean	1995-07-25 12:00:00	3.654225e+05
min	1967-01-07 00:00:00	1.620000e+05
25%	1981-04-16 06:00:00	2.910000e+05
50%	1995-07-25 12:00:00	3.420000e+05
75%	2009-11-01 18:00:00	3.990000e+05
max	2024-02-10 00:00:00	6.137000e+06
std	NaN	2.418473e+05


```
In [440]: plt.figure(figsize=(12, 6))
plt.plot(df.index, df['ICSA'])
plt.title('ICSA time series plot')
plt.xlabel('Year')
plt.ylabel('ICSA')
plt.grid(True)
plt.show()
```



```
In [441]: # Outlier Detection
plt.figure(figsize=(12, 6))
sns.boxplot(df['ICSA'])
plt.title('Box Plot for Outlier Detection')
plt.show()
```




```
In [442]: # Distribution of ICSA values
plt.figure(figsize=(12, 6))
sns.histplot(df['ICSA'], kde=True)
plt.title('Distribution of ICSA Values')
plt.xlabel('ICSA')
plt.ylabel('Frequency')
plt.show()
```



```
In [443]: # removing outliers
Q1 = df['ICSA'].quantile(0.25)
Q3 = df['ICSA'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
print(lower_bound, upper_bound)
df_filt = df[(df['ICSA'] >= lower_bound) & (df['ICSA'] <= upper_bound)]
```

```
129000.0 561000.0
```


```
In [444]: plt.figure(figsize=(12, 6))
plt.plot(df_filt.index, df_filt['ICSA'])
plt.title('ICSA time series plot after outlier removal')
plt.xlabel('Year')
plt.ylabel('ICSA')
plt.grid(True)
plt.show()
```



```
In [445]: # Boxpot after outlier removal
plt.figure(figsize=(12, 6))
sns.boxplot(df_filt['ICSA'])
plt.title('Box Plot after Outlier removal')
plt.show()
```



```
In [446]: # Distribution of ICSA values after outlier removal
plt.figure(figsize=(12, 6))
sns.histplot(df_filt['ICSA'], kde=True)
plt.title('Distribution of ICSA Values after outlier removal')
plt.xlabel('ICSA')
plt.ylabel('Frequency')
plt.show()
```




```
In [447]: # Checking for stationarity and making the Series Stationary if necessary
result = adfuller(df_filt['ICSA'])
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
print('Critical Values:')
for key, value in result[4].items():
    print(f'    {key}: {value}')
```

```
ADF Statistic: -3.382131
p-value: 0.011583
Critical Values:
  1%: -3.4326595388027648
  5%: -2.8625603948435945
 10%: -2.5673131867249634
```

As p-value is less than 0.05 and absolute of ADF Statistic (-3.38) < absolute of Critical Value (5%) (-2.862), time series is stationary.

```
In [448]: plot_acf(df_filt['ICSA'])
          plot_pacf(df_filt['ICSA'])
          plt.show()
```



```
In [449]: df.dropna(inplace=True)
```

```
In [450]: from sklearn.model_selection import train_test_split
          train, test = train_test_split(df_filt['ICSA'], test_size=0.2, random_state=45)
```

```
In [451]: model = ARIMA(train, order=(2, 0, 2))
          results = model.fit()
```

C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: An unsupported index was provided and will be ignored when e.g. forecasting.

self._init_dates(dates, freq)

C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: An unsupported index was provided and will be ignored when e.g. forecasting.

self._init_dates(dates, freq)

C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: An unsupported index was provided and will be ignored when e.g. forecasting.

self._init_dates(dates, freq)

```
In [452]: print(results.summary())
```

SARIMAX Results

```
=====
Dep. Variable:          ICSA      No. Observations:      2270
Model:                ARIMA(2, 0, 2)  Log Likelihood      -28808.227
Date:                 Thu, 22 Feb 2024  AIC                57628.454
Time:                 06:35:34    BIC                57662.820
Sample:              0      HQIC                57640.992
                        - 2270
Covariance Type:      opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
const      3.375e+05    1775.535     190.079     0.000     3.34e+05     3.41e+05
ar.L1       0.0492       0.427       0.115     0.908     -0.787       0.886
ar.L2       0.7103       0.309       2.297     0.022       0.104       1.316
ma.L1      -0.0232       0.428      -0.054     0.957     -0.862       0.816
ma.L2      -0.7189       0.317      -2.267     0.023     -1.340      -0.097
sigma2     6.184e+09       0.002    2.71e+12     0.000     6.18e+09     6.18e+09
=====
=
Ljung-Box (L1) (Q):          0.55   Jarque-Bera (JB):          16.8
1
Prob(Q):          0.46   Prob(JB):          0.0
0
Heteroskedasticity (H):      0.89   Skew:          0.2
0
Prob(H) (two-sided):      0.11   Kurtosis:          2.8
4
=====
=
```

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
[2] Covariance matrix is singular or near-singular, with condition number 4.6e+27.
Standard errors may be unstable.
```



```
In [453]: for i in range(1, 4):
            df_filt[f'ICSA_lag{i}'] = df_filt['ICSA'].shift(i)

df_filt.dropna(inplace=True)
X = df_filt[['ICSA_lag1', 'ICSA_lag2', 'ICSA_lag3']]
Y = df_filt['ICSA']

X = sm.add_constant(X)

Xtrain, Xtest, Ytrain, Ytest = train_test_split(X, Y, test_size=0.2, random_state=42)

model = sm.OLS(Ytrain, Xtrain)
results = model.fit()
print(results.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          ICSA      R-squared:                0.958
Model:                  OLS      Adj. R-squared:           0.958
Method:                 Least Squares      F-statistic:        1.716e+04
Date:                  Thu, 22 Feb 2024      Prob (F-statistic):    0.00
Time:                  06:35:35      Log-Likelihood:       -25211.
No. Observations:      2268      AIC:                  5.043e+04
Df Residuals:          2264      BIC:                  5.045e+04
Df Model:               3
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	4862.2616	1506.593	3.227	0.001	1907.813	7816.710
ICSA_lag1	0.6993	0.020	34.372	0.000	0.659	0.739
ICSA_lag2	0.1031	0.026	4.020	0.000	0.053	0.153
ICSA_lag3	0.1834	0.021	8.635	0.000	0.142	0.225

```
=====
Omnibus:                 1463.714      Durbin-Watson:           1.971
Prob(Omnibus):            0.000      Jarque-Bera (JB):        85826.171
Skew:                     2.341      Prob(JB):                0.00
Kurtosis:                 32.771      Cond. No.                2.64e+06
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.64e+06. This might indicate that there are strong multicollinearity or other numerical problems.

C:\Users\Akhila Markunda\AppData\Local\Temp\ipykernel_15852\3266025240.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_filt[f'ICSA_lag{i}'] = df_filt['ICSA'].shift(i)
```

C:\Users\Akhila Markunda\AppData\Local\Temp\ipykernel_15852\3266025240.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_filt[f'ICSA_lag{i}'] = df_filt['ICSA'].shift(i)
```

C:\Users\Akhila Markunda\AppData\Local\Temp\ipykernel_15852\3266025240.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_filt[f'ICSA_lag{i}'] = df_filt['ICSA'].shift(i)
```

C:\Users\Akhila Markunda\AppData\Local\Temp\ipykernel_15852\3266025240.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_filt.dropna(inplace=True)
```

```
In [454]: Ypred = results.predict(Xtest)
print("The forecasted values are")
print(Ypred)
```

The forecasted values are

```
2666    227364.099630
1817    480688.972285
1464    329879.210121
261     274443.463596
487     378117.046020
```

...

```
1753    304574.522611
85      199371.191380
104     216353.720673
2362    378954.875396
1190    357709.960042
```

Length: 567, dtype: float64

```
In [455]: url3 = "https://fred.stlouisfed.org/graph/fredgraph.csv?id=IC4WSA"
```

```
In [456]: response2 = requests.get(url3)
df_2 = pd.read_csv(io.StringIO(response2.text))
print(df_2.head(10))
```

	DATE	IC4WSA
0	1967-01-28	209000
1	1967-02-04	211000
2	1967-02-11	216500
3	1967-02-18	219500
4	1967-02-25	229000
5	1967-03-04	252500
6	1967-03-11	255500
7	1967-03-18	259500
8	1967-03-25	260750
9	1967-04-01	248000

```
In [457]: # Check for missing values
missing_values = df_2.isnull().sum()
# Check for duplicates
duplicates = df_2.duplicated().sum()

# Handling missing values (if any, here we just print them)
print(f"Missing values:\n{missing_values}")
print(f"Duplicate values:\n{duplicates}")
```


```
Missing values:
DATE      0
IC4WSA    0
dtype: int64
Duplicate values:
0
```

Here there are no missing or duplicate values in IC4WSA data.

```
In [458]: print(f"Statistical Summary of IC4WSA data:")
print(df_2.describe())
```

```
Statistical Summary of IC4WSA data:
          IC4WSA
count  2.977000e+03
mean   3.655887e+05
std    2.289940e+05
min    1.790000e+05
25%    2.910000e+05
50%    3.415000e+05
75%    3.990000e+05
max    5.288250e+06
```

```
In [459]: plt.figure(figsize=(12, 6))
plt.plot(df_2.index, df_2['IC4WSA'])
plt.title('IC4WSA time series plot')
plt.xlabel('Year')
plt.ylabel('IC4WSA')
plt.grid(True)
plt.show()
```




```
In [460]: df_2['DATE'] = pd.to_datetime(df_2['DATE'])
df_merge = pd.merge(df_filt, df_2, on='DATE', how='inner')
print(df_2.dtypes)
```


```
DATE      datetime64[ns]
IC4WSA      int64
dtype: object
```

```
In [461]: corr = df_merge[['ICSA', 'IC4WSA']].corr()
```

```
In [462]: plt.figure(figsize=(12, 6))
plt.scatter(df_merge['IC4WSA'], df_merge['ICSA'])
plt.xlabel('IC4WSA')
plt.ylabel('ICSA')
plt.title('Scatter Plot of ICNSA vs IC4WSA')
plt.show()
```



```
In [463]: period_s = 12
decomp_result = seasonal_decompose(df_merge['ICSA'], model='additive', period=period_s)
decomp_result.plot()
plt.show()
```



```
In [469]: X = df_merge['IC4WSA']
Y = df_merge['ICSA']

Xtrain, Xtest, Ytrain, Ytest = train_test_split(X, Y, test_size=0.2, random_state=42)
```

```
In [472]: model = ARIMA(Xtrain,order=(2, 0, 2))
result = model.fit()
print(result.summary())
```

```
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473:
ValueWarning: An unsupported index was provided and will be ignored when e.g. fore
casting.
    self._init_dates(dates, freq)
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473:
ValueWarning: An unsupported index was provided and will be ignored when e.g. fore
casting.
    self._init_dates(dates, freq)
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473:
ValueWarning: An unsupported index was provided and will be ignored when e.g. fore
casting.
    self._init_dates(dates, freq)
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:9
66: UserWarning: Non-stationary starting autoregressive parameters found. Using ze
ros as starting parameters.
    warn('Non-stationary starting autoregressive parameters'
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:9
78: UserWarning: Non-invertible starting MA parameters found. Using zeros as start
ing parameters.
    warn('Non-invertible starting MA parameters found.'
```

SARIMAX Results

```

=====
Dep. Variable:          IC4WSA      No. Observations:          2268
Model:                ARIMA(2, 0, 2)  Log Likelihood          -28799.071
Date:                 Thu, 22 Feb 2024  AIC              57610.142
Time:                 06:37:49         BIC              57644.502
Sample:               0              HQIC             57622.679
                             - 2268

```

```

Covariance Type:      opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	3.374e+05	1852.790	182.087	0.000	3.34e+05	3.41e+05
ar.L1	-0.1285	0.197	-0.651	0.515	-0.515	0.258
ar.L2	0.8486	0.192	4.423	0.000	0.473	1.225
ma.L1	0.1467	0.206	0.712	0.477	-0.257	0.551
ma.L2	-0.8359	0.201	-4.149	0.000	-1.231	-0.441
sigma2	6.28e+09	0.003	2.1e+12	0.000	6.28e+09	6.28e+09

```

=====
=

```

```

Ljung-Box (L1) (Q):          0.00   Jarque-Bera (JB):          21.4
3
Prob(Q):                     0.98   Prob(JB):              0.0
0
Heteroskedasticity (H):      0.90   Skew:                  0.2
3
Prob(H) (two-sided):         0.13   Kurtosis:              2.9
0

```

```

=====
=

```

Warnings:

```

[1] Covariance matrix calculated using the outer product of gradients (complex-step).
[2] Covariance matrix is singular or near-singular, with condition number 1.8e+27.
Standard errors may be unstable.

```



```
In [473]: f = result.get_forecast(1)
f_v = f.predicted_mean
c_i = f.conf_int()
print("Forecast values:")
print(f_v)
print("\nConfidence intervals:")
print(c_i)
```

```
Forecast values:
2268    334217.179433
dtype: float64
```

```
Confidence intervals:
          lower IC4WSA  upper IC4WSA
2268  178899.052795  489535.30607
```

```
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836:
ValueWarning: No supported index is available. Prediction results will be given wi
th an integer index beginning at `start`.
```

```
    return get_prediction_index(
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836:
FutureWarning: No supported index is available. In the next version, calling this
method in a model without a supported index will result in an exception.
    return get_prediction_index(
```

```
In [ ]:
```